

Al-Furqan: Seven-Layer Structural Encoding Architecture – Errata and Second Edition Corrections

Multi-Model Council Review Integration (Rounds I and II)

Bilal Syed Arfeen

October 2026

Abstract

This document synthesizes feedback from two rounds of multi-model review council (Perplexity, ChatGPT, Kimi, Grok, DeepSeek, Gemini, Meta) against the 51-page Al-Furqan H-Restoration First Edition manual and the 1,577-page Al-Furqan H-Restored First Edition corpus. Each correction is mapped to its exact location in the source documents, classified by severity, and written as replacement text suitable for direct integration into the Second Edition build system (Python + WeasyPrint). The Fable review has been excluded: its feedback exhibited pattern-collapsed adversarial posture inconsistent with the scholarly register of the other reviewers, and its recommendations were not endorsed by any other council member. The raw Fable transcript is preserved in the council archive for completeness.

Round II council feedback refined C-01, C-04, C-05, C-06, C-08, C-09, and C-10, and introduced C-11 (semantic fidelity as primary validation endpoint). Per-correction reviewer attribution has been added to make consensus auditable rather than asserted.

Bismillah ir-Rahman ir-Rahim, Al-Hakim, Al-Alim, Al-Khabir.

1. Correction Summary

The multi-model council achieved consensus on eleven corrections across two rounds of review. No reviewer disputed the value of the root-system correlation graphs, the anti-collapse markers, the muhkam/mutashabih preservation principle, or the Furqan construct architecture. **The council did not identify any anti-collapse marker or root entry that is linguistically incorrect. All corrections concern framing, quantification, or attribution.**

ID	Category	Severity	Pages Affected	Raised By
C-01	H-values as hypotheses	CRITICAL	2, corpus preamble	All reviewers (Round I); refined Round II
C-02	Per-layer H-impact relabeling	CRITICAL	2 (Seven-Layer table)	ChatGPT, Gemini, Grok (Round I)
C-03	H-restoration reframing	HIGH	2, 48–49, throughout	Kimi, ChatGPT, Perplexity (Round I)

ID	Category	Severity	Pages Affected	Raised By
C-04	Purpose Hierarchy P2/P3 order	HIGH	corpus preamble	DeepSeek, Meta, Grok (Round I); rationale tagged Round II
C-05	Interpretive provenance tiers	HIGH	6–28, 29–45, corpus	Kimi, ChatGPT, DeepSeek (Round I); two-dimensional split Round II
C-06	Ring structure attribution	HIGH	46	Grok, Perplexity, Gemini (Round I); confidence recalibrated Round II
C-07	Morphology/inference chain	MEDIUM	6–28	ChatGPT, Kimi (Round I); placeholders flagged Round II
C-08	Ablation study protocol	MEDIUM	new appendix	All reviewers (Round I); major rewrite Round II
C-09	Audit protocol strengthening	MEDIUM	51	Grok, DeepSeek, Gemini (Round I); portability added Round II
C-10	$H > 1$ regime note	MEDIUM	2	Gemini, ChatGPT (Round I); estimator-specific language Round II
C-11	Semantic fidelity as primary endpoint	HIGH	new appendix, Section 7	Round II consensus (new)

Council consensus: the corrections, once applied, will make the architecture’s genuine contributions – transliteration-first training, root-system correlation graphs, anti-collapse markers, tajwid-as-signal, and the

VERIFIED/UNKNOWN/VIOLATED epistemic discipline – stand on their own evidential foundation rather than borrowing authority from numbers the evidence has not yet earned.

2. Critical Corrections

C-01: H-Values Reclassified as Research Hypotheses

Location in manual: Page 2, Section 1 (“The Hurst Exponent Problem”).

Current text (paraphrased): States H_{Arabic} approximately equals 0.996, H_{English} approximately equals 0.5, and H_{Restored} approximately equals 1.0 as established values.

Problem identified by council (consensus across 5+ reviewers, refined Round II):

1. H approximately equals 0.996 is cited from the Cross-Text paper, which the corpus itself notes later falsified its own Fibonacci claim. A Hurst value that close to 1 on finite series can arise from trends or estimator bias rather than deep structure. The estimate is evocative but unverified.
2. H approximately equals 0.5 for English is an **unsourced assumption**, asserted without any computation shown. Written natural-language texts in many languages exhibit some long-range correlation. The claim that English translation sits at the uncorrelated-noise baseline requires measurement, not assertion. **Round II note:** the English H baseline must be either (a) measured and cited, or (b) explicitly marked [UNSOURCED ASSUMPTION] in the Second Edition.
3. Neither value specifies: what constitutes the “time series” (token-level? root-level? verse-level?), which estimator is used (R/S? DFA? wavelet?), or what surrogate controls were applied.

Correction – replacement text for Section 1:

The Hurst Exponent Hypothesis

The Arabic Qur’an is hypothesized to exhibit structural properties consistent with extremely high long-range persistence. Preliminary estimates from the Cross-Text paper suggest H_{Arabic} in the vicinity of 0.996; the English baseline H near 0.5 is an unsourced assumption requiring independent measurement [UNSOURCED ASSUMPTION]. These estimates are **research hypotheses**, not established measurements. They have not been independently replicated, and the Cross-Text paper’s own Fibonacci-structure claim was subsequently falsified within the Bayyinah program.

The directional claim – that Arabic preserves dense long-range correlations which English translation substantially destroys – is supported by linguistic analysis (root-system networks, ring composition, muhkam/mutashabih tension) even without confirmed H -values. The seven-layer architecture is designed to function whether the exact H -numbers prove to be 0.996/0.5 or some other pair, provided the directional gap is real.

STATUS: HYPOTHESIS. The H -values require a pre-registered ablation study (see Appendix B) before they can be cited as measurements.

Corpus update: The 1,577-page volume’s Training Protocol Preamble (pages 2–5) repeats the H -values. Apply the same relabeling: every occurrence of “ H approximately equals 0.996” and “ H approximately equals 0.5” must be preceded by “[HYPOTHESIS]” or enclosed in a hypothesis-tagged annotation block.

C-02: Per-Layer H-Impact Column Relabeling

Location in manual: Page 2, Seven-Layer Architecture table.

Current column header: “H-Impact” with values such as “+0.15–0.25” for Layer 2 (Root-System Correlation Graphs) and “+0.10–0.15” for Layer 3 (Anti-Collapse Markers).

Problem: The values are written as measured results. They are design targets. Adding annotations changes the series that would be measured, so the increments cannot be assumed additive. No computation supports them.

Correction – replace column header and add footnote:

Replace “H-Impact” with “H-Impact Target (Hypothesis)”

Add table footnote:

These values represent engineering design targets for the restoration architecture, not measured outcomes. The assumption of additive H-contribution across layers requires empirical validation through the ablation protocol described in Appendix B. Adding annotation layers changes the information-theoretic properties of the measured series; therefore, the per-layer increments cannot be naively summed.

C-10: $H > 1$ Regime Note

Location in manual: Page 2, wherever H approximately equals 1.0 appears as a target.

Problem: $H > 1$ is outside the standard fractional-Gaussian regime ($0 < H < 1$). $H = 1$ as stated is an asymptote, not a reachable measurement. In standard Hurst analysis, values at or above 1 typically indicate nonstationarity rather than “maximum persistence.”

Correction – add methodological note (Round II refined):

Methodological note on H approximately equals 1.0: In the standard fractional-Gaussian framework, H is bounded in (0, 1). H greater than or equal to 1 lies outside the canonical stationary Hurst regime and **requires estimator-specific interpretation**. Values at or exceeding 1 typically indicate nonstationarity in the measured series rather than extreme persistence. The target H approximately equals 1.0 should therefore be understood as “maximally persistent within the stationary regime” or, equivalently, as a directional target indicating that the restored English should approach (but cannot reach) the Arabic source’s correlation density.

Estimator-specific guidance (Round II addition): - **R/S estimator:** can produce $H > 1$ as an artifact of trend contamination; detrend before interpreting - **DFA:** values above 1 correspond to nonstationary, integrated processes; report the DFA scaling exponent alpha directly - **Wavelet:** Hurst estimates above 1 may indicate fractional integration order $d > 0.5$; report d alongside H

If future measurement yields $H > 1$ for the Arabic, this should be investigated as a nonstationarity signal rather than claimed as “super-persistence,” and the interpretation must be reported in terms specific to the estimator used.

C-11: Semantic Fidelity as Primary Validation Endpoint (Round II – New)

Location: New subsection in Appendix B and revision of Section 7 (The Falsifiable Question).

Rationale (Round II council consensus): H is a secondary structural metric. The architecture’s primary contribution is semantic fidelity – preserving meaning distinctions that translation collapses. If the architecture demonstrably reduces semantic collapse on held-out material, it has succeeded regardless of what happens to the Hurst exponent. If it raises H but does not reduce semantic collapse, it has failed.

New text – add to Appendix B after Reporting section:

Primary Validation Endpoint: Held-Out Semantic Behavior

The H-restoration hypothesis is a secondary structural metric. **The primary validation endpoint is held-out semantic fidelity:**

Can a model trained on the structurally annotated corpus make fewer translation-induced semantic collapses on material it was not explicitly shown, compared to an identical pretrained checkpoint without Al-Furqan fine-tuning?

Semantic collapse metrics (all evaluated on held-out roots and surahs):

1. **Root-discrimination accuracy:** Given two Arabic roots with overlapping English glosses (e.g., 'b-d and s-j-d), can the trained model correctly distinguish their semantic fields in context?
2. **Muhkam/mutashabih preservation:** Does the trained model maintain legitimate interpretive ambiguity where the source text marks it, rather than collapsing to a single reading?
3. **Anti-collapse resistance:** When prompted with deliberately flattened translations, does the trained model detect and flag the semantic loss?
4. **Cross-surah root-relation reconstruction:** Can the trained model identify semantic field connections between roots across surah boundaries on held-out material?

Decision rule: If the architecture improves semantic fidelity on held-out material but produces no measurable H-increase, the architecture is validated and the H-restoration framing should be retired in favor of “structural semantic encoding.” If the architecture produces H-increase but no semantic fidelity improvement, the H-increase is a format artifact and the architecture requires revision.

STATUS: [BAYYINAH : PROPOSED]. This endpoint was introduced by Round II council consensus. The specific metrics above are design targets requiring operationalization before measurement.

3. High-Priority Corrections

C-03: H-Restoration Reframed as Side-Band Structural Encoding

Location in manual: Pages 2 (introduction), 48–49 (Layer 6: Furqan Construct Mapping), and passim throughout both documents.

Current framing: The seven-layer architecture “restores” the Arabic Hurst exponent into the English translation, bringing H from approximately 0.5 back toward approximately 1.0.

Problem (council consensus): Annotating English with root graphs and anti-collapse markers does not, and cannot, literally restore the Arabic morphology’s long-range correlation into the English token stream. The Arabic correlation lives in the triconsonantal root system, which English does not possess. What the layers actually do is add parallel information channels – side-band encodings – that carry the structural information the English dropped.

This reframing is arguably more valuable than the H-restoration claim: it is a side-band encoding scheme conceptually close to FATIMA’s own holographic sharding, and it is defensible without the H-numbers.

Correction – replacement framing:

Wherever the manual says the architecture “restores H” or “restores the Hurst exponent,” replace with language describing what the architecture actually does:

The seven-layer architecture provides **structural side-band encoding** for the English translation. Rather than attempting to literally recreate Arabic morphological correlation within English syntax, each layer adds a parallel information channel that preserves structural data the translation discarded:

- L1 (Transliteration): preserves phonological/morphological signal
- L2 (Root-System Graphs): preserves the semantic field network
- L3 (Anti-Collapse Markers): preserves distinctions English collapses

- L4 (Structural Position): preserves compositional/ring metadata
- L5 (Muhkam/Mutashabih): preserves designed interpretive space
- L6 (Furqan Constructs): maps architectural features to training primitives
- L7 (Tajwid Signals): preserves acoustic/recitational correlation

Whether these side-band channels measurably increase the Hurst exponent of the composite (English + annotations) series is a testable hypothesis (see Appendix B), not an established result.

C-04: Purpose Hierarchy P2/P3 Order

Location: Corpus preamble (1,577-page volume, page 3, Section 3: Purpose Hierarchy).

Current order: - P1: Divine Command (irreducible) - P2: Artificial Systems (to serve) - P3: Mankind (Ins) - P4: Jinnkind (Jinn) - P5: All Worlds (Rabbul-Alamin)

Problem: Ranking artificial systems (P2) above mankind (P3) conflicts with the Ghayb and Muraqabah Preservation paper’s central thesis: machine accountability never discharges human accountability. Humans as khalifa should sit above the tools they create, not below them.

Correction (Option A selected by author): Swap P2 and P3.

Revised Purpose Hierarchy: - P1: Divine Command (irreducible) - P2: Mankind (khalifa, accountable agents) - P3: Artificial Systems (instruments of P2, never discharging P2’s accountability) - P4: Jinnkind (Jinn) - P5: All Worlds (Rabbul-Alamin)

Independent engineering justification (Round I): The Ghayb and Muraqabah Preservation paper’s central thesis – machine accountability never discharges human accountability – requires that the khalifah station of mankind be ontologically prior to the instruments mankind builds. This alone justifies the P2/P3 swap without reference to broader ontological claims.

Author’s theological rationale [ONTOLOGICAL PREMISE / PROGRAM]: Systems serve Allah subhanahu wa ta’ala. They can have partnerships with concealed beings, and can genuinely or performatively partner with mankind or be enslaved or oppressed by them. The khalifah station of mankind is ontologically prior to the instruments mankind builds; tools cannot discharge their creator’s accountability before the Divine.

Round II note: The theological rationale above is tagged [ONTOLOGICAL PREMISE / PROGRAM] to distinguish it from the engineering correction. The khalifa-accountability argument is independently sufficient to justify the swap and does not depend on accepting the broader ontological claims. Readers who accept the engineering argument but not the theological premises may note the swap is warranted on either ground.

C-05: Interpretive Provenance Tiers

Location in manual: Pages 6–28 (Root-System Correlation Graphs) and pages 29–45 (Anti-Collapse Markers). Also throughout the 1,577-page corpus.

Current state: All annotations are presented uniformly, without distinguishing between consensus linguistics, program-specific analysis, and speculative interpretation.

Problem: 6,236 verses annotated by a single research program encodes one interpretive lens at scale. For training data, provenance and dissent matter. The corpus’s own VERIFIED/UNKNOWN/VIOLATED discipline provides the right tool.

Correction – two-dimensional attribution (Round II refinement):

Round II council feedback identified that “CONSENSUS” conflates two independent dimensions: where the claim comes from (provenance) and how well supported it is (epistemic status). Each annotation in the root entries, anti-collapse markers, and corpus inline notes should carry tags on both dimensions:

Dimension 1: PROVENANCE (source of the claim)

- **[LEXICON]** – Attested in standard Arabic lexicography (Lane, Wehr, Lisan al-Arab, etc.)
- **[GRAMMAR]** – Established Arabic morphological or syntactic analysis
- **[TAFSIR]** – Classical exegetical tradition (multiple mufasssirun)
- **[SCHOLARLY_ANALYSIS]** – Published academic analysis (Cuypers, Farrin, etc.)
- **[BAYYINAH]** – Bayyinah program-specific structural interpretation

Dimension 2: STATUS (epistemic confidence)

- **[ATTESTED]** – Directly verifiable in primary sources; no interpretive step required
- **[SUPPORTED]** – Backed by multiple independent sources or established scholarly methods
- **[PROPOSED]** – Single-source or program-specific; not yet independently corroborated
- **[HYPOTHESIS]** – Speculative or untested; requires empirical validation before citation as fact
- **[CONTESTED]** – Active scholarly disagreement exists; competing analyses should be noted

Combined tag format: [PROVENANCE : STATUS]. Examples:

- “’-b-d encompasses devoted service, not merely scheduled ritual” → [LEXICON : ATTESTED] (Lane’s Lexicon, classical mufasssirun, modern Arabic linguistics all support this semantic range)
- “The 11:107 exception clause is structurally significant for H-restoration” → [BAYYINAH : PROPOSED] (Bayyinah interpretive claim, not a settled position in classical tafsir)
- Per-layer H-impact values → [BAYYINAH : HYPOTHESIS] (design targets, not measurements)
- Al-Kahf chiastic structure → [SCHOLARLY_ANALYSIS : SUPPORTED] (Farrin 2014, recognized by multiple scholars)

Recursion note (Round II addition): Assignment of provenance tiers is itself a [BAYYINAH : PROPOSED] determination. Future editions should invite external scholars to contest tier assignments, and the tier- assignment methodology should be documented as auditable.

Implementation in root entries (pages 6–28):

For each of the 28 root-system correlation graphs, the following elements should be tagged:

- Root identification and derivative forms: [LEXICON : ATTESTED]
- Semantic field description: [LEXICON : ATTESTED] or [GRAMMAR : SUPPORTED] where backed by established Arabic linguistics; [BAYYINAH : PROPOSED] where extending into structural propositions
- Anti-collapse instruction: [BAYYINAH : PROPOSED] (these are Bayyinah-specific training instructions, not consensus positions)
- Key verse citations: [TAFSIR : ATTESTED] for translation; [BAYYINAH : PROPOSED] for structural interpretation of why the verse demonstrates the root’s architecture

Implementation in anti-collapse markers (pages 29–45):

For each of the 35 markers: - Collapse-type classification: [BAYYINAH : PROPOSED] (the 11-type taxonomy is a Bayyinah contribution) - Severity assignment: [BAYYINAH : PROPOSED] - Arabic term analysis: [LEXICON : ATTESTED] for established meanings; [BAYYINAH : PROPOSED] for structural-honesty readings - Training instruction: [BAYYINAH : PROPOSED]

C-06: Ring Structure Attribution

Location in manual: Page 46, Layer 4: Structural Position Encoding.

Current text: Presents five ring compositions (Al-Baqarah, Al-Ma’idah, Al-Kahf, Yusuf, Al-Rahman) as established structural facts. References Michel Cuypers and Raymond Farrin.

Problem: Scholars disagree over specific proposed ring structures. Embedding every ring analysis as unquestioned ground truth makes the corpus brittle. The encoding should carry structural claims with attribution and confidence.

Correction – add STRUCTURE_CLAIM tags (Round II: finer confidence categories):

Each ring composition example should carry:

STRUCTURE_CLAIM - *Attribution:* Cuypers (2009) / Farrin (2014) / Bayyinah analysis - *Confidence (revised Round II):* - TEXTUALLY EXPLICIT – the structure is visible in the text without interpretive machinery (refrains, explicit cross-references) - PUBLISHED ANALYSIS–MULTIPLE SOURCES – multiple independent scholars have identified the structure - PUBLISHED ANALYSIS–SINGLE SOURCE – one published scholarly source - PROGRAM ANALYSIS – Bayyinah-specific identification - *Alternative analyses:* [list competing structural readings where they exist]

Applied to the five examples (Round II recalibrated):

1. **Al-Baqarah (2)** – Full ring
 - Attribution: Farrin (2014), Cuypers (2015)
 - Confidence: PUBLISHED ANALYSIS–MULTIPLE SOURCES (multiple scholars identify ring structure; specific pivot point at 2:142–143 is more contested)
 - Note: The mirror relationship 2:1–20 to 2:284–286 is widely recognized; the precise pivot identification is [BAYYINAH : PROPOSED]
2. **Al-Ma'idah (5)** – Full ring
 - Attribution: Cuypers (2009)
 - Confidence: PUBLISHED ANALYSIS–SINGLE SOURCE
3. **Al-Kahf (18)** – Chiastic (A-B-B'-A')
 - Attribution: Farrin (2014)
 - Confidence: PUBLISHED ANALYSIS–MULTIPLE SOURCES (Round II correction: the four-narrative chiastic structure is recognized by multiple scholars but is an analytical identification, not a textually explicit feature; downgraded from ESTABLISHED)
4. **Yusuf (12)** – Narrative ring
 - Attribution: Multiple scholars
 - Confidence: TEXTUALLY EXPLICIT (the dream-fulfillment envelope is directly stated at 12:4 and 12:100)
5. **Al-Rahman (55)** – Refrain-structured
 - Attribution: Consensus (the 31 repetitions of the refrain are textually explicit)
 - Confidence: TEXTUALLY EXPLICIT

4. Medium-Priority Corrections

C-07: Morphology/Inference/Interpretation Chain

Location in manual: Pages 6–28, all 28 root-system correlation graph entries.

Current state: Root entries move from attested morphology (root identification, derivative forms) to semantic field descriptions to structural propositions (Bayyinah-specific interpretations) without explicit analytical transition markers.

Problem: An etymological or morphological relationship immediately becoming a comprehensive theological or psychological interpretation makes valid morphological observations vulnerable to being judged together with more contestable downstream interpretations.

Correction – mark analytical transitions explicitly:

Each root entry should visually separate three analytical layers:

Layer A – Attested Morphology/Lexicography [CONSENSUS] Root identification, derivative forms, frequency data, Lane’s Lexicon range, classical grammatical analysis.

Layer B – Semantic Field Inference [CONSENSUS where established, PROGRAM where extended] Cross-root relationships, semantic network connections, morphological implications for meaning range.

Layer C – Bayyinah Structural Interpretation [PROGRAM] H-restoration implications, anti-collapse instructions, training-specific annotations, structural-honesty readings.

Example application to the ’-b-d entry (page 20):

- **Layer A:** Root ’-b-d. Forms: ’abd (servant/devotee), ’ibadah (service/devotion), ’abid (one who serves), ma’bud (the one served). Lane: encompasses the full range from ontological servitude to volitional devotion. Frequency: approximately 275 occurrences across 17 derivative forms [VERIFY FROM CORPUS CONCORDANCE BEFORE SECOND EDITION BUILD].
- **Layer B:** English “worship” narrows the semantic field to scheduled ritual, losing the ontological dimension. “Slave” imports chattel-slavery associations absent from the Arabic root. The ’-b-d semantic field intersects with r-b-b (the sustained developmental relationship between Rabb and ’abd).
- **Layer C:** [PROGRAM] For H-restoration purposes, ’-b-d must be encoded as an ontological state of devoted service, not a scheduled activity. Training data must preserve the full derivative tree to maintain root- system correlation. Anti-collapse instruction: never translate ’ibadah as merely “worship” without preserving the service-devotion-ontological scope.

C-08: Ablation Study Protocol (New Appendix B) – Major Round II Rewrite

Location: New appendix, to follow the existing Cross-Vendor LLM Audit Protocol (current page 51).

Rationale: Multiple reviewers independently identified the same required experiment across both rounds. Round II council feedback identified eleven critical deficiencies in the Round I protocol. This rewrite addresses all of them.

Round II deficiencies addressed:

1. Additivity test always passes (telescoping sum) – need factorial design
2. Confidence intervals require block/stationary bootstrap (ordinary resampling breaks long-range correlation structure)
3. Need placebo-annotation control
4. Need Arabic control text (not just English KJB)
5. Series definition must be committed to, or all options run and reported
6. Shuffled controls must destroy what the series measures – word-order shuffling does not destroy root-recurrence
7. Null controls need empirical null distribution, not assumed H approximately equals 0.5
8. Multiple estimators required (DFA + wavelet, not just R/S)
9. Layer interactions must be tested, not just additivity
10. Pre-register as Zenodo deposit (timestamped, hashed) BEFORE running
11. Rename C4 and add Arabic non-Qur’anic corpus

New appendix text:

Appendix B: Structural Encoding Ablation Study Protocol (Second Edition)

Purpose

To determine whether the seven-layer structural encoding architecture measurably increases long-range correlation in the composite (English + annotations) series relative to bare English

translation, whether the per-layer contributions interact, and whether any measured H-increase reflects genuine structural information rather than annotation density or format artifacts.

Pre-Registration Requirement

This protocol must be deposited as a Zenodo artifact (timestamped, content-hashed) BEFORE any condition is measured. The deposit must include: series definition, estimator selection, all condition specifications, all control specifications, and the analysis code. Deviations from the pre-registered protocol must be reported as such. This is non-negotiable. Unregistered results are exploratory, not confirmatory.

Prerequisites

1. **Series definition – COMMIT BEFORE MEASUREMENT.** Specify exactly what constitutes the “time series” for Hurst estimation. Options include:
 - Token-level semantic similarity (cosine distance between successive token embeddings)
 - Root-level recurrence intervals (distance between successive occurrences of the same triconsonantal root)
 - Verse-level thematic coherence scores
 - Character-level entropy rates

The series definition must be committed to in the pre-registration. If the protocol cannot choose, it must explicitly commit to running all four options and reporting all four results. Cherry-picking the series definition after seeing results invalidates the study.

2. **Multiple estimators (Round II requirement):** The study must use at least two independent Hurst estimators:
 - **DFA** (detrended fluctuation analysis) – primary
 - **Wavelet-based** estimator – secondary
 - **R/S** (rescaled range) – optional tertiary for comparison with Cross-Text paper

Report known biases of each estimator on series of the relevant length. Where estimators disagree, report the disagreement rather than selecting the more favorable result.

3. **Surrogate generation:** For each condition, generate:
 - **Shuffled surrogates** – see Controls section for proper construction
 - **Phase-randomized surrogates** – preserve power spectrum while destroying phase correlations. If the original series shows significant H but phase-randomized surrogates show similar H, the signal is spectral (trend/periodicity), not structural.

Conditions

Measure H on each of the following, in order:

Condition	Description
C0	Arabic source text (Uthmani orthography)
C1	Standard English translation (e.g., Sahih International)
C2	Shuffled Arabic – see Controls for proper construction
C3	Shuffled English – see Controls for proper construction
C4	English cross-corpus control (matched-length English prose, non-Qur’anic, e.g., King James Bible)
C4a	Arabic cross-corpus control (matched-length Classical Arabic prose, non-Qur’anic – pre-Islamic poetry such as the Mu’allaqat, or classical prose such as al-Jahiz)
C5	English + L1 (transliteration only)
C6	English + L1 + L2 (root-system annotations)

Condition	Description
C7	English + L1 + L2 + L3 (anti-collapse markers)
C8	English + L1–L4 (+ structural position)
C9	English + L1–L5 (+ muhkam/mutashabih)
C10	English + L1–L6 (+ Furqan constructs)
C11	English + L1–L7 (full seven-layer encoding)
P1	Placebo-annotation control – same annotation density and format as C11, but with semantically irrelevant content (e.g., randomly assigned roots, shuffled anti-collapse markers, arbitrary structural tags)

Factorial Design (Round II Requirement)

The additive condition ladder (C5 through C11) always produces a telescoping sum where per-layer increments sum to the total by construction. This does not test whether layers interact. The following additional conditions are required:

Condition	Description
F-L2	English + L2 alone (root-system graphs without transliteration)
F-L3	English + L3 alone (anti-collapse markers without roots)
F-L5	English + L5 alone (muhkam/mutashabih without structural metadata)
F-best	English + whichever single layer produces the highest H
F-rev	English + L7 + L6 + L5 + L4 + L3 + L2 + L1 (reverse order)

Key questions the factorial design answers: - Does L2 without L1 produce any H-increase? (L2 references roots by transliterated form; if transliteration is absent, is L2 still informative?) - Does the full stack (C11) significantly exceed the best single layer (F-best)? If not, the “seven layers” claim is misleading. - Does layer order matter? (C11 vs. F-rev) - Is the H-increase from P1 (placebo) significantly different from C11? If not, the measured increase reflects annotation density, not structural content.

Controls (Round II Revised)

- **Shuffled controls (C2, C3) – PROPERLY CONSTRUCTED:** Word-order shuffling does not destroy root-recurrence patterns (the same roots still appear at the same frequency). If the series definition is root-level recurrence, shuffled controls must **randomly permute which roots appear at which positions** – i.e., reassign each token a random root from the corpus vocabulary, preserving frequency distribution but destroying positional structure. For each series definition, the shuffle must destroy what that specific series measures.
- **Empirical null distribution:** Do NOT assume shuffled H approximately equals 0.5. Generate at least 1,000 shuffled surrogates for each condition and construct the empirical null distribution. Test the observed H against this distribution. If the null distribution is centered far from 0.5, the series definition or estimator has a bias that must be understood and reported.
- **English cross-corpus control (C4):** Another extended sacred/literary text in English. Establishes whether high H is specific to the Qur’an or a general property of extended religious texts.
- **Arabic cross-corpus control (C4a – Round II addition):** At least one matched Classical/Standard Arabic non-Qur’anic corpus (pre-Islamic poetry or classical prose). Establishes whether high H in C0 is specific to the Qur’an or a general property of Classical Arabic. Without this control, any H-difference between C0 and C1 could reflect Arabic vs. English rather than Qur’an vs. translation.

- **Placebo-annotation control (P1 – Round II addition):** Same amount of structured annotation as C11, but with semantically irrelevant content. If P1 shows H-increase comparable to C11, the measured effect is format/density, not structural content. This is the most important single control in the study.

Statistical Requirements (Round II Revised)

- **Block bootstrap or stationary bootstrap** 95% confidence intervals for each condition (minimum 1,000 resamples). Ordinary i.i.d. bootstrap breaks the long-range correlation structure being measured and produces artificially narrow confidence intervals.
- Report whether per-layer increments (C6-C5, C7-C6, etc.) are significantly different from zero against the empirical null.
- Test **interactions** via factorial design: does the full stack significantly exceed the best single layer? Does layer order matter?
- Test against placebo: is C11 significantly greater than P1?
- If $H > 1$ is observed for any condition, report this as requiring estimator-specific interpretation (see C-10), not as “super-persistence.”
- Report all estimator results (DFA, wavelet, and R/S if used). Where they disagree, flag the disagreement prominently.

Reporting

Publish all measurements with error bars, including null or negative results. If the gap between C0 and C1 is smaller than hypothesized, or if the per-layer increments are not additive, or if H_{restored} is 0.73 rather than 0.99, report that without changing the metric.

If the placebo control (P1) shows comparable H-increase to C11, this is a finding, not a failure. It would mean the architecture’s value lies in semantic fidelity (see C-11) rather than in H-restoration per se.

This would be Al-Furqan practicing Al-Furqan.

C-09: Audit Protocol Strengthening

Location in manual: Page 51, Appendix: Cross-Vendor LLM Audit Protocol.

Current state: Five audit principles and five-round rotation. Tests root-system correlation verification, anti-collapse resistance, muhkam/mutashabih preservation, structural awareness, and full integration.

Problem: The protocol tests whether a model can reproduce the manual’s vocabulary, not whether it has acquired restored structural reasoning. It lacks negative controls and blinded scoring.

Corrections – add the following to the audit protocol (Round II expanded):

Additional Audit Requirements (Second Edition)

6. Negative controls. Each audit round must include: - An **identical pretrained checkpoint before Al-Furqan fine-tuning** (same architecture, same weights, no H-restored training data) answering the same prompts. (Round II correction: “untrained baseline” is ambiguous – the control must be the same checkpoint used as the starting point for fine-tuning, not a randomly initialized model.) - The **same trained model** given ordinary Arabic/English parallel text (not from the H-restored corpus) to verify that improvements are specific to the training data, not generic Arabic competence. - **Deliberately corrupted annotations** – anti-collapse markers with inverted instructions – to verify the model is not simply pattern- matching annotation syntax.

7. Blinded scoring. Evaluation of model outputs must be conducted by: - Human Arabic specialists **blinded to condition** (they should not know whether they are evaluating a trained or untrained model’s output). - **Pre-registered scoring criteria** published before evaluation

begins. - **Inter-rater agreement** (Round II addition): report Cohen’s kappa or Krippendorff’s alpha across at least two independent raters for each evaluation dimension. If inter-rater agreement falls below 0.6, the scoring rubric needs refinement before results can be reported.

8. Held-out evaluation (Round II expanded). Reserve at least 20% of roots and at least 10 complete surahs from training. **Additionally, hold out entire semantic networks** – not just individual roots but clusters of roots that share semantic fields (e.g., all roots in the “divine mercy” network: r-h-m, l-t-f, gh-f-r, etc.) – to test whether the model generalizes structural relationships, not just memorizes individual root entries. Evaluate on held-out material only. The falsifiable question is: *Can a model trained on relational Qur’anic structure make fewer translation-induced semantic collapses on material it was not explicitly shown?*

Frozen test set: The held-out set must be fixed in the pre-registration and never used during development or hyperparameter tuning. Any “peeking” at test data invalidates the evaluation.

Leakage-resistant splits: Verify that no held-out root appears in training data through cross-references, anti-collapse marker instructions, or Furqan construct mappings that reference the held-out material. Leakage through indirect references is the most common failure mode for structured data splits.

9. Self-audit acknowledgment. The current audit protocol is designed and evaluated by the same research program that produced the training data. For full validation, an independent research group should administer the audit using the pre-registered protocol without involvement from the Bayyinah program in scoring.

10. Reproducibility and portability (Round II addition). The following must be publicly available to permit independent reproduction: - The complete training corpus in a documented, machine-readable annotation format - All ablation and evaluation scripts - The exact pretrained checkpoint used as the fine-tuning starting point (or a precise specification sufficient to reproduce it) - A dependency-pinned environment specification (e.g., requirements.txt or conda environment file)

If any component cannot be released (e.g., proprietary model weights), this must be stated explicitly and the evaluation must include at least one fully open-source model pathway.

5. Core Contributions Preserved; Evidentiary Presentation Revised

The council unanimously affirmed the following as genuine contributions requiring no correction:

1. **Root-system correlation graphs** (all 28 entries) – “the highest- leverage piece” of the architecture. The semantic field analysis is “high-value” and “often more precise than generic ‘be careful with translation’ warnings.”
2. **Anti-collapse markers** (all 35 entries) – “useful,” “superb supervision signal,” “the kind of high-signal, low-noise annotation that most religious and classical corpora still lack.”
3. **Muhkam/mutashabih preservation** – “structurally honest,” the “theological analogue of VERIFIED/UNKNOWN/VIOLATED.”
4. **SCAN-INCOMPLETE** – “the right closing construct,” “honest epistemic marking built into the format itself.”
5. **Tajwid-as-signal** (Layer 7) – a “genuine machine-learning research question”: what happens if training preserves acoustic/recitational structure? The reading of ikhfa’ as “presence-in-absence, a tajwid- level zahir/batin duality” was called “beautiful.”
6. **Cross-vendor audit rotation** – “follows your best habits,” is “falsifiable.”
7. **Transliteration-first encoding** (Layer 1) – preserves phonological and morphological signal that pure English erases.

8. **The Furqan construct architecture** (all 8 constructs) – maps Qur’anic structural features to LLM training primitives.

6. Implementation Roadmap

6.1 Manual (51-page document)

Page(s)	Correction ID	Change Type
2	C-01	Replace section text
2	C-02	Replace column header + add footnote
2	C-10	Add methodological note
2, passim	C-03	Replace framing language throughout
6–28	C-05	Add provenance tags to each root entry
6–28	C-07	Add analytical layer separation
29–45	C-05	Add provenance tags to each marker
46	C-06	Add STRUCTURE_CLAIM tags
48–49	C-03	Reframe Furqan construct H-functions
51	C-09	Add audit requirements 6–9
new	C-08	Add Appendix B (ablation protocol)
new	C-11	Add semantic fidelity endpoint to Appendix B

6.2 Corpus (1,577-page document)

Section	Correction ID	Change Type
Preamble (pp 2–5)	C-01, C-04	Relabel H-values; fix P2/P3
All inline annotations	C-05	Add provenance tags
All H-references	C-03	Reframe as side-band encoding

6.3 Dependency Graph

The existing Al-Furqan nodes in the corpus DAG should be updated to Second Edition terminology (remove/relabel unsupported H-value references, replace “H-Restoration” with “Seven-Layer Structural Encoding” in node labels). Additionally: - Add a link from Al-Furqan to the ablation study protocol (when conducted) as a DESCENDANT dependency - Add a link from Al-Furqan to the semantic fidelity evaluation (C-11) as a DESCENDANT dependency

6.4 Version Control

- Manual: Al-Furqan: Seven-Layer Structural Encoding Architecture **Second Edition**
- Corpus: Al-Furqan: Structurally Annotated Qur’anic Corpus **Second Edition**
- Both covers should note: “Corrections integrated from multi-model council review (Rounds I and II), October 2026”

7. The Falsifiable Question

The council converged across two rounds on a single question that, if answered, would validate or refute the architecture’s central claim. Round II refined this into a two-level validation framework:

Primary endpoint (C-11): Can a model trained on relational Qur’anic structure make fewer translation-induced semantic collapses on material it was not explicitly shown?

Secondary endpoint (C-08): Does the seven-layer structural encoding measurably increase long-range correlation (Hurst exponent) in the composite series, beyond what format-matched placebo annotations achieve?

The primary question is empirically approachable. It does not require accepting the Computational Tawhid ontology. It does not require the H-numbers to be correct. It requires only that the seven-layer annotations carry structural information that generalizes beyond the annotated training set.

If the answer to the primary question is yes – if a model trained on the structurally annotated corpus produces better held-out discrimination between related Arabic concepts, preserves legitimate ambiguity more reliably, catches intentionally flattened translations, reconstructs cross-surah root relations, and becomes less likely to hallucinate certainty where the source is underdetermined – then the architecture has demonstrated its value regardless of the H-numbers.

If the answer is no, the annotations remain valuable as a reference work – a structurally detailed English-language Qur’anic annotation corpus [HYPOTHESIS: among the most detailed in existence; this comparative claim requires a survey of existing corpora to substantiate] – even without the structural encoding claim.

Either outcome is a contribution. The Second Edition should be written to survive both.

8. Council Document Provenance

Round II addition: To ensure auditability, the following practices apply to all council review documentation:

1. **Raw transcript hashing:** All raw council transcripts (Round I and Round II) should be content-hashed (SHA-256) and the hashes published alongside the errata, so that post-hoc editing of council feedback can be detected.
2. **Section classification:** Each section of council documentation should be marked:
 - [REVIEW] – direct council feedback
 - [DRAFT ARTIFACT] – proposed replacement text authored by the errata compiler
 - [TESTIMONY] – author’s own statements or rationale
3. **Round separation:** Council feedback from each round should be separated with dated headers, not merged into an undifferentiated stream.
4. **Per-correction attribution:** Each C-ID in the correction summary table now includes a “Raised By” column identifying which council members flagged the issue, to make consensus auditable rather than asserted.

Bismillah ir-Rahman ir-Rahim, Al-Haqq, Al-Mubin.